

Continuous Variables

Biology 683

Lecture 5

Heath Blackmon

Last week

1. Give a couple of examples of pseudoreplication.
2. Give an example of an experiment from your own field where you could use a binomial or chi square tests

Today

- 1) Comparing means
 - a) One sample t-test
 - b) Two sample t-test
 - c) Paired value t-test
- 2) Comparing Variances
 - a) Levene's test
- 3) Assumptions
- 4) Transformations
- 5) Non-parametric tests
 - a) Sign test
 - b) Mann-Whitey / Wilcoxon Rank
 - c) Permutation test

Continuous Data (more or less)

Many datasets are technically categorical but can be treated as continuous.

- count data**

 - chromosomes**

 - offspring**

 - species**

 - gene copy number**

t tests

- Based on the t -distribution, which is similar to a normal distribution, except it takes into account the fact that we don't know the population standard deviation with certainty. Published by Gosset in 1908 under the pseudonym Student because he was employed by Guinness brewery.
- Since we only know the sample standard deviation, instead of the population standard deviation, we have some additional uncertainty regarding the possible values of the mean
- Consequently, the t -distribution is wider than the normal distribution, and the t -distribution actually yields a more exact estimate of the confidence interval for the mean

BIOMETRIKA.

THE PROBABLE ERROR OF A MEAN.


By STUDENT.

Introduction.

ANY experiment may be regarded as forming an individual of a "population" of experiments which might be performed under the same conditions. A series of experiments is a sample drawn from this population.



T distribution



Power increases dramatically
for the first few samples

One-sample t-test

The one sample T-test compares a sample to a hypothesis


Assumptions:

- The variable is normally distributed in the population

Implementation

But this is a proportion

```
x <- c(.45, .43, .31,  
       .54, .47, .39)  
  
t.test(x, mu=.5)
```



One Sample t-test


```
data: x  
t = -2.1579, df = 5, p-value = 0.08341  
alternative hypothesis: true mean is not equal to 0.5  
95 percent confidence interval:  
 0.3502649 0.5130684  
sample estimates:  
mean of x  
0.4316667
```

$$t = \frac{(\bar{Y} - \mu)}{SEM}$$

One-sample t-test

Arcsine is the standard transformation for proportion data. Lets see how our answer changes.

```
x <- c(.45, .43, .31,  
       .54, .47, .39)  
tx <- asin(sqrt(x))  
t.test(tx, mu=asin(sqrt(.5)))
```



$$t = \frac{(\bar{Y} - \mu)}{SEM}$$

One Sample t-test

```
data: tx  
t = -2.1444, df = 5, p-value = 0.08484  
alternative hypothesis: true mean is not equal to 0.7853982  
95 percent confidence interval:  
 0.6329179 0.7991801  
sample estimates:  
mean of x  
 0.716049
```


If you transform don't forget to back transform

One Sample t-test

```
data: tx
t = -2.1444, df = 5, p-value = 0.08484
alternative hypothesis: true mean is not equal to 0.7853982
95 percent confidence interval:
 0.6329179 0.7991801
sample estimates:
mean of x
 0.716049
```

```
x <- c(.45, .43, .31,
        .54, .47, .39)
tx <- asin(sqrt(x))
res <- t.test(tx, mu=asin(sqrt(.5)))
sin(res$conf.int)^2
[1] 0.3498723 0.5137801
```

Two-sample t-test


Many times we will want to compare two populations strains or treatments and see if their means differ.

Implementation

```
# population 1 measures
x <- c(.45, .43, .31, .37,
      .47, .25, .39, .22)

# population 2 measures
y <- c(.42, .49, .45, .49,
      .59, .37, .57, .53)

t.test(x, y)
```



Welch Two Sample t-test

```
data: x and y
t = -3.0297, df = 13.377, p-value = 0.009405
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.21815501 -0.03684499
sample estimates:
mean of x mean of y
 0.36125  0.48875
```

$$t = \frac{(\bar{Y}_a - \bar{Y}_b)}{\sqrt{\frac{s_a^2}{n_a} + \frac{s_b^2}{n_b}}}$$

Paired-sample t-test


In some experiments you measure individuals before and after a manipulation. These values can be analyzed with the paired-sample t-test.

$$t = \frac{\bar{d}}{SE_d} \quad \bar{d} = \frac{\sum_{i=1}^n y_{ai} - y_{bi}}{n}$$

Implementation

```
# first measure
x <- c(.24, .16, .14, .33, .04)
# second measure
y <- c(.52, .41, .10, .47, .39)


t.test(x, y, paired = T)
```




Paired t-test

```
data: x and y
t = -2.8817, df = 4, p-value = 0.04493
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.384839093 -0.007160907
sample estimates:
mean of the differences
 -0.196
```

Confidence Intervals and t-tests




If two 95% confidence intervals don't overlap then the values are significantly different at the $\alpha = 0.05$ level.



If there is overlap in the CI then you can't conclude anything.

Comparing variances


- Some experiments might focus on differences in variance. Think about how sexual selection works with regard to number offspring.



- Two most common tests are the F-test and Levine's test. F-test assumes that the variable is normally distributed in the population and is sensitive to violations of this assumption.
- For this reason we will focus on Levine's test which assumes the variables have roughly symmetrical distributions and is furthermore quite robust to violations of this assumption.

Levene's test

- Levene's test works by calculating a central value for each group and then comparing the deviations of individuals in each group from these to calculate a W statistic.
- Compared to a F distribution



Levene's test


```
library(car)
data("InsectSprays")
leveneTest(y = InsectSprays$count,
           group = InsectSprays$spray,
           center = "median")
```

Levene's Test for Homogeneity of Variance (center = "median")

	Df	F value	Pr(>F)
group	5	3.8214	0.004223 **
	66		

Signif. codes:

0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1



Assumptions

- **One-sample T-test**
 - The variable is normally distributed in the population
- **Two-sample T-test**
 - The variable is normally distributed in the population
 - unequal variance ok
- **Paired-sample T-test**
 - Differences are normally distributed
- **Levine's test**
 - Distributions are roughly symmetrical

Assumptions

- My data doesn't look normal!
- What do I do?




Assumptions

- Many tests are quite robust to minor violations of assumptions.
- Tests of normality (Shapiro-Wilks) are not terribly useful.
 - Small sample size have little power
 - Large sample size reject with even minor deviations from normality
- Often most important to simply visualize distributions
 - histograms
 - qqplots


```
shapiro.test(x)
```

histograms

Histogram of rnorm(50)




Histogram of rnorm(50)



You don't need or even expect a perfect bell curve from a finite sample of a population.


histograms




You should be worried if the histogram suggests a different distribution

qqplots

`qqnorm(x)`



normal



Too many extreme values

The time when you should be most hesitant is when your biological knowledge tells you that you don't expect a variable to have a normal distribution!

Data really isn't close to normal

First, be sure that the statistical test actually requires normality as an assumption – this constraint isn't universal (e.g., Levene's test merely requires symmetry)

Second, try to **transform** the variable to make it roughly normal

Third, as a last resort, use a non-parametric test

Transformations

Log transformation

- Take the natural log of every observation
- Especially for right-skewed data.

```
log(x)  
# if zeros are present  
log(x + 1)
```

Arcsine transformation

- Take the arcsine of the square root of each observation
- Used mainly for proportions

```
asin(sqrt(x))
```

Square-root transformation

- Add 0.5 and take the square root
- Solves similar problems to the log transformation

```
sqrt(x + .5)
```

For other transformations, see p. 381 in Whitlock and Schluter

Don't P-hack

- Try different transformations to see which improve normality or homoscedasticity not to see which one gives the lowest p-value.
- If you use a transformation that is not standard in your field or for the data type you will often need to justify it in the text of your paper or to the reviewers of your work
- Best case scenario is when you can take two routes and show the same result.
 - *The number of genes in the P53 network that showed signs of selection were positively correlated with the maximum life span of taxa ($B = 2.5$, $R^2 = .26$, $p = .024$). Analysis of square root transformed gene count data produced qualitatively similar results.*
 - *A Welch's two sample t-test showed that the two populations had significantly different mean body size. However despite log transformation of body size both populations failed a Shapiro-Wilks test of normality. Therefore we also performed a Mann-Whitney U-test which also showed that population A was significantly larger than population B.*

Non-parametric tests

1. Non-parametric tests still have assumptions but fewer.
2. You lose a lot of power when you switch from parametric to non-parametric.

Mann-Whitey / Wilcoxon Rank

Replaces two sample t-test

- Tests the hypothesis based upon a ranked list
- The tests are done on the ranks rather than the actual values
- Implementation: `wilcox.test(x, y)`
- Assumptions:
 - The distribution of measurements is symmetrical
 - Distributions of the two variables must have the same shape
 - A **significant result implies that the two variables have different distributions** (possibly due to different means but also possibly due to variances or skewness)

Doing a Mann-Whitey / Wilcoxon Rank in R

```
x <- c(45, 43, 46, 47,  
       47, 25, 39, 22)  
y <- c(42, 49, 45, 49,  
       59, 37, 57, 53)  
wilcox.test(x, y)
```

Wilcoxon rank sum test with continuity
correction

```
data: x and y  
W = 14.5, p-value = 0.07357
```

```
> t.test(x,y)
```

Welch Two Sample t-test

```
data: x and y  
t = -2.1694, df = 12.883, p-value = 0.04937  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
-19.21901046 -0.03098954  
sample estimates:  
mean of x mean of y  
39.250 48.875
```

Sign test

Replaces one-sample or paired-sample t-test

- Convert the data to a score (+ or -)
 - One-sample: whether the value is above or below the hypothesized value.
 - Paired-sample: whether the second value is higher or lower
- Under the null hypothesis, the number of +'s should equal the number of -'s, so this test boils down to a binomial test
- Obviously, it's throwing away a lot of information
- Implementation: `binom.test(x=4, n=14, p=.5)`

Doing sign test in R

- Implementation:

```
x <- c(45, 43, 31, 37,
       47, 25, 39, 22)
y <- c(42, 49, 45, 49,
       59, 37, 57, 53)
z <- x>y
[1] TRUE FALSE FALSE FALSE FALSE FALSE FALSE
binom.test(x=sum(z), n=length(z), p=.5)
```

Exact binomial test

```
data:  sum(z) and length(z)
number of successes = 1, number of trials = 8,
p-value = 0.07031
alternative hypothesis: true probability of success is not equal to 0.5
95 percent confidence interval:
 0.003159724 0.526509671
sample estimates:
probability of success
      0.125
```

Test	P-value
T-test	0.049
Wilcox rank test	0.074
Sign test	0.125

Permutation test

- Makes almost no assumptions about the underlying data and allow very specific questions to be asked.
- Under the null hypothesis, the two groups are drawn from identical distributions (they are from identical populations)
- By randomizing the observations between groups, an expected distribution under the null hypothesis can be generated
- If we are interested in differences between the mean, then we might choose $\text{mean 1} - \text{mean 2}$ as our test statistic

Permutation test

Step 1: Create a permuted set of data in which the values of the response variable are randomly reordered

Step 2: Calculate the test statistic from the randomly reordered groups


Step 3: Repeat Steps 1 and 2 thousands of times

Step 4: Compare the actual value of the test statistic from the original data to the distribution of values from the permutation steps

Permutation test

```
crickets <- read.csv("crickets.csv")
obs.vals <- tapply(X = crickets$timeToMating,
                  INDEX = crickets$feedingStatus,
                  FUN = mean)
obs.d <- obs.vals[2]-obs.vals[1]
perm <- 10000
perm.d <- vector()
for(i in 1:perm){
  newX <- sample(crickets$timeToMating, replace = F)
  perm.vals <- tapply(X = newX,
                    INDEX = crickets$feedingStatus,
                    FUN = mean)
  perm.d[i] <- perm.vals[2] - perm.vals[1]
}
plot(density(perm.d))
abline(v = obs.d, col = "red")
sum(perm.d < obs.d) / perm * 2

t.test(crickets$timeToMating~crickets$feedingStatus)
```



Permutation test (vers. 2)

```
# READ IN THE DATA
dat <- read.csv("crickets.csv")

# this data has an odd distribtuion lets
# use it to illustrate the permutation test
hist(dat$timeToMating[12:24])

# this will hold our samples of the null distribution
null.dist <- vector()


# for our permutation number we will use iter
iter <- 10000

# here we create the statistical distribution
# by taking the same size sample as we did
# in our experiment and calculating the mean
for(i in 1:iter){
  null.dist[i] <- mean(sample(dat$timeToMating, size=11))
}

ybar <- mean(dat$timeToMating[1:11])

pval <- sum(null.dist <= ybar)/10000
pval
hist(null.dist)
abline(v=ybar, col="red")

t.test(x = dat$timeToMating[1:11],
       y= dat$timeToMating[12:24],
       alternative = "less")
```



P-value from permutation test = 0.0647

P-value from T-test = 0.058

For Thursday

Read chapter WS 10-13

Bring laptop to class!

Heath Blackmon

BSBW 309

coleoguy@gmail.com

@coleoguy